



# Exam AZ-203

**Developing Solutions  
for Microsoft Azure**

# Agenda

- Storage Account
- Blob Storage
- Content Delivery Network
- Table Storage

# Blob Types

©Smruti Ranjan



Block Blob




Append Blob




Page Blob

# Blob Structure

©Smruti Ranjan





<http://mystorageaccount.blob.core.windows.net/container/Image1.jpg>

<http://mystorageaccount.blob.core.windows.net/container?comp=list>

# Access Levels



## Private

Authenticated request



## Blob

Anonymous read access for blobs but cannot list all the blobs inside container.



## Container

All container and blob data can be read by anonymous request. Clients can enumerate blobs within the container by anonymous request but cannot enumerate containers within the storage account.

# Security

©Smruti Ranjan



Account key



SAAS token

# Features



Handles  
concurrency



Soft Delete  
feature



Create  
Snapshots



# Demo

- Create a container and Blobs
- Showcase access level
- Soft Delete
- .Net SDK



# Concurrency



**1. Optimistic concurrency** – An application performing an update will as part of its update verify if the data has changed since the application last read that data. For example, if two users viewing a wiki page make an update to the same page then the wiki platform must ensure that the second update does not overwrite the first update – and that both users understand whether their update was successful or not. This strategy is most often used in web applications.



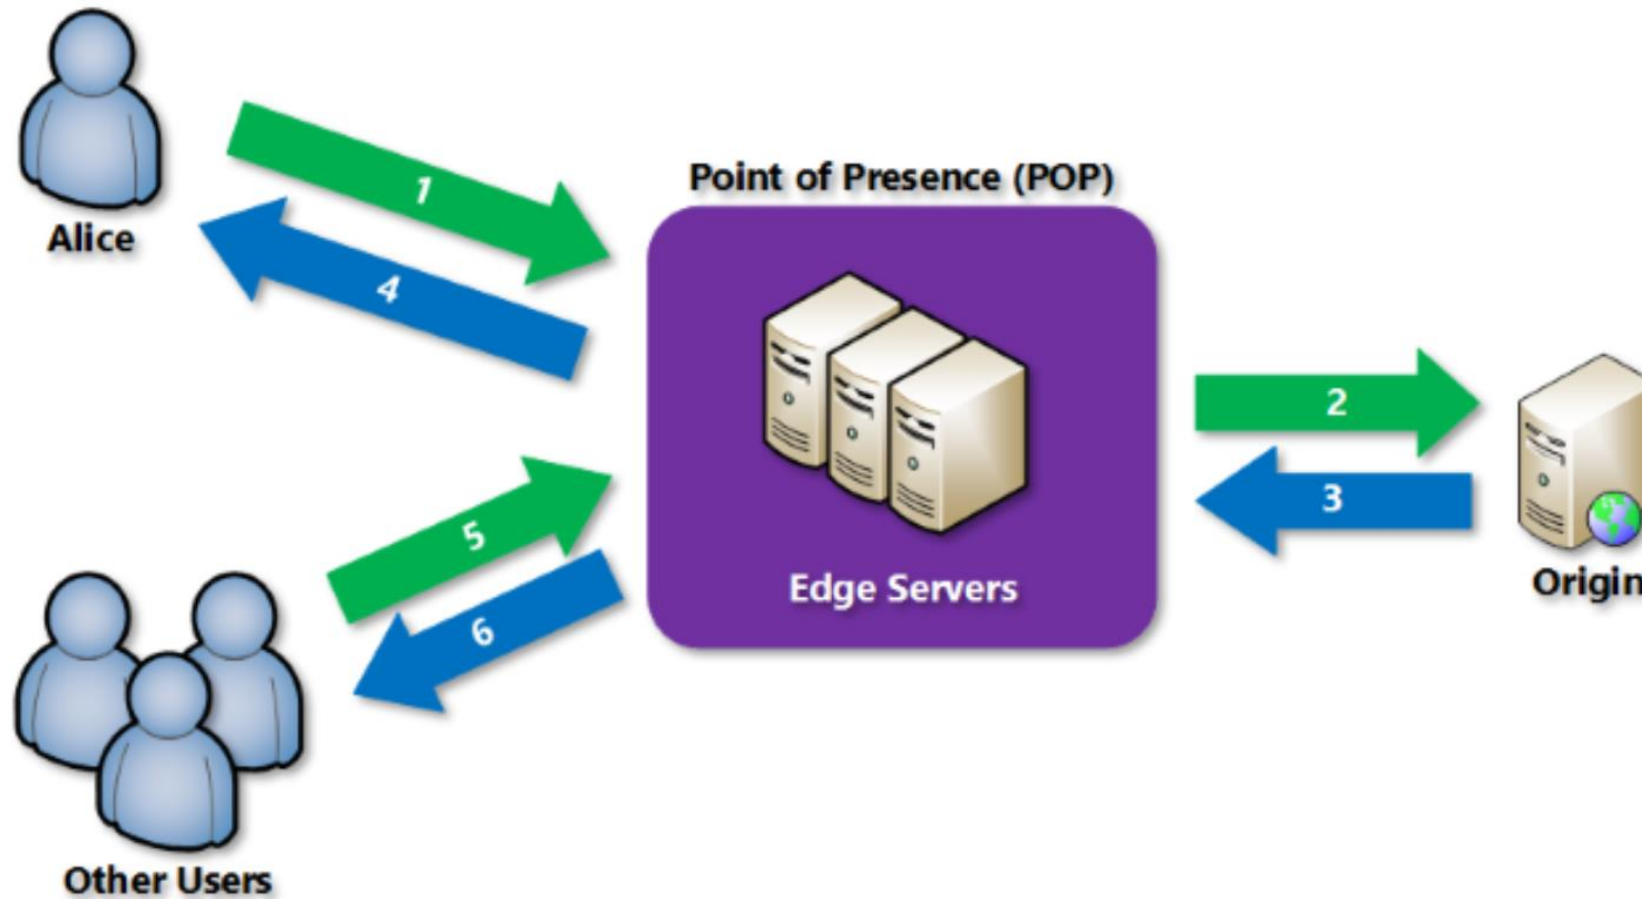
**2. Pessimistic concurrency** – An application looking to perform an update will take a lock on an object preventing other users from updating the data until the lock is released. For example, in a master/slave data replication scenario where only the master will perform updates the master will typically hold an exclusive lock for an extended period of time on the data to ensure no one else can update it.



**3. Last writer wins** – An approach that allows any update operations to proceed without verifying if any other application has updated the data since the application first read the data. This strategy (or lack of a formal strategy) is usually used where data is partitioned in such a way that there is no likelihood that multiple users will access the same data


# Azure CDN Overview

©Smruti Ranjan



# Cache Rules

- **Ignore query strings:** Default mode. In this mode, the CDN point-of-presence (POP) node passes the query strings from the requestor to the origin server on the first request and caches the asset. All subsequent requests for the asset that are served from the POP ignore the query strings until the cached asset expires.
- **Bypass caching for query strings:** In this mode, requests with query strings are not cached at the CDN POP node. The POP node retrieves the asset directly from the origin server and passes it to the requestor with each request.
- **Cache every unique URL:** In this mode, each request with a unique URL, including the query string, is treated as a unique asset with its own cache. For example, the response from the origin server for a request for `example.ashx?q=test1` is cached at the POP node and returned for subsequent caches with the same query string. A request for `example.ashx?q=test2` is cached as a separate asset with its own time-to-live setting.




# Azure Table Storage

©Smruti Ranjan



NoSQL



Uses JSON to  
serialize the data



Flexible schema  
changes



Scalable and highly  
available

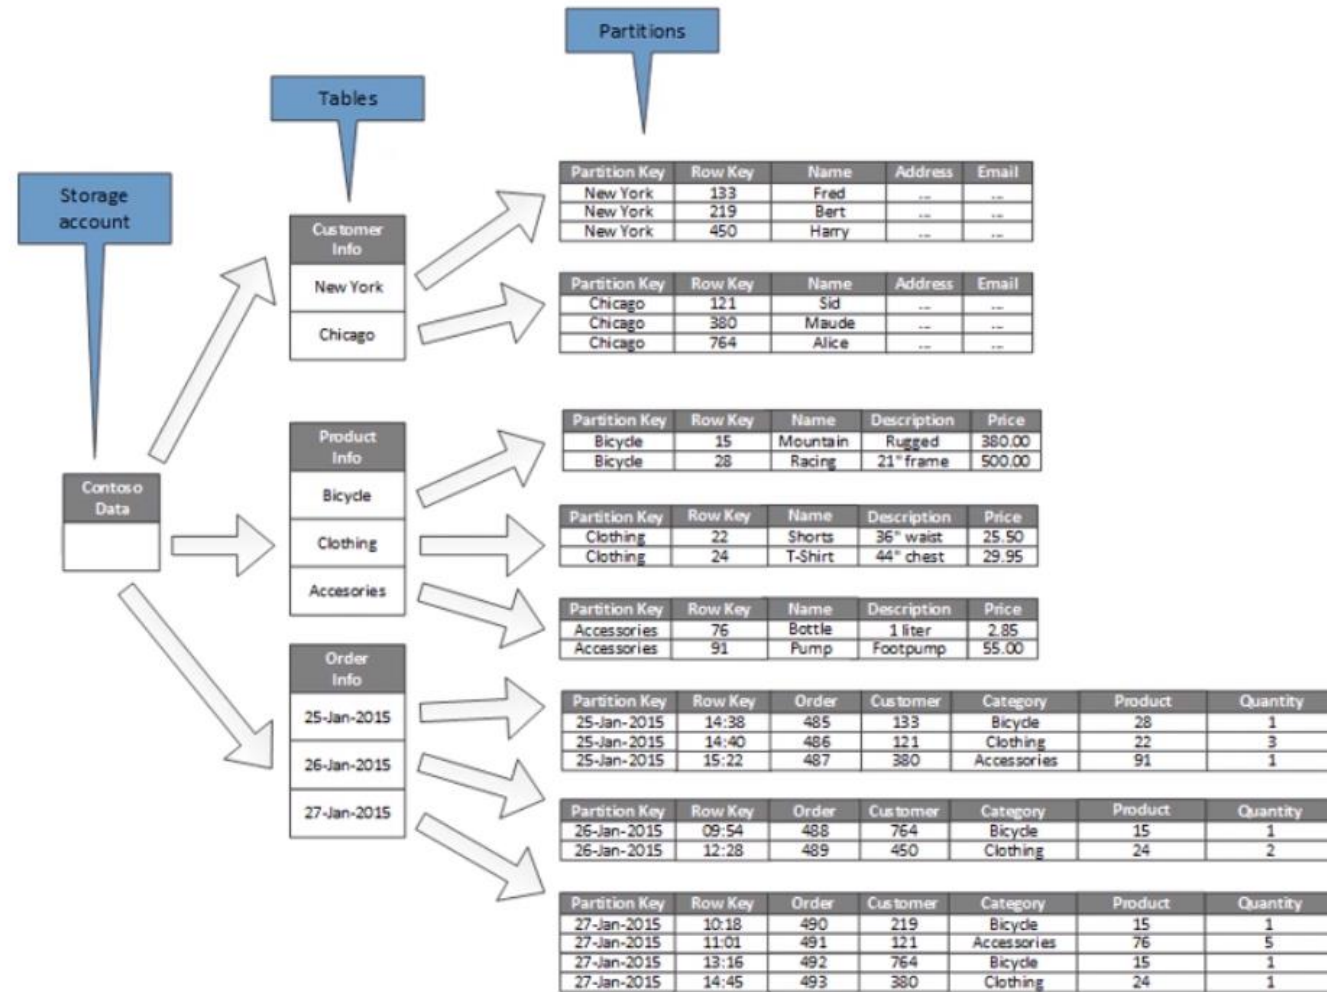


Single region with  
optional secondary



read-only region

# Partitioning schemes



# Demo

- Table Storage Rest API

